

Parent Selection and Diversification in Genetic Programming

Thomas Helmuth
Computer Science
Washington and Lee Univ
Lexington, Virginia
helmuth@wlu.edu

Nicholas Freitag McPhee
Div of Sci and Math
Univ of MN, Morris
Morris, MN 56267
mcphee@morris.umn.edu

Lee Spector
Cognitive Science
Hampshire College
Amherst, MA
lspector@hampshire.edu

ABSTRACT

More things!

Keywords

lexicase selection, hyperselection, PushGP, other stuff

1. INTRODUCTION

I bet we start here!

Lexicase selection [4] is nifty, eh?

2. EXPERIMENTAL DESIGN

Previous work has shown that using lexicase selection results in higher population error diversity than tournament selection across a variety of problems [1, 3]. These examples examined the diversity of entire GP runs, each starting with a different initial population and random number seed.

Here we examine the effects of these parent selection methods on population diversity starting from specific population conditions besides a random initial population. In particular, we want to see how each method changes diversity in populations that occur naturally during an evolutionary run.

In order to produce the populations on which to experiment, we started GP runs and let them continue until they met certain stopping conditions; we then stored those populations and later conducted multiple trials with different random number seeds starting with those stored populations. We used three different stopping conditions in order to generate naturally occurring populations with interesting properties:

1. In a run using lexicase selection, we stopped if the population error diversity was greater than 0.9. This results in very diverse populations, allowing us to observe whether evolution is able to maintain such high diversity in the following generations.
2. In a run using tournament selection, we stopped if the population error diversity was less than 0.15. These

populations allow us to see if methods promote diversification starting from such undiverse populations. They also allow us to see if methods perform differently on a population produced by tournament selection versus one produced by lexicase selection.

3. As described above, we were initially motivated here by observations of runs using lexicase selection undergoing major drops in diversity following hyperselection events, where one or a few individuals in the population received the majority of the parent selections in a generation. We had anecdotally noticed rapid diversity recovery following these events, but not examined them systematically.

In this condition, we stopped a run using lexicase selection when the error diversity reached a level at least 0.25 less than it had been at some point in the previous 10 generations. This allowed us to detect populations that had recently undergone large drops in diversity. We do not know for sure whether those drops are related to hyperselection events, but we expect that they are.

In all three conditions, we only considered populations occurring after generation 10 in order to give evolution a chance to settle down after the extreme shifts that can happen at the beginning of a run.

In each trial, we continued running GP on a stored population for 20 generations and recorded the population error diversity. For each parent selection setting (lexicase and tournament selections), we conducted 100 trials with different random number seeds from each stored population.

We conducted these tests on two problems taken from a recent program synthesis benchmark suite [2]. The first problem, Replace Space With Newline (RSWN), searches for a program that takes as input a string and both prints the string after replacing all of the spaces in the input with newline characters and functionally returns the number of non-whitespace characters in the string. Previous examinations of error vector diversity on the RSWN problem indicate that lexicase selection maintains significantly higher diversity than tournament selection, which across 100 runs never had a median diversity higher than 0.25 [1].

The second problem, Double Letters, asks for a program that takes a string as input and prints the string after doubling every alphabetic character and tripling every exclamation point. All other characters should be printed once. As with the RSWN problem, lexicase selection consistently achieves high diversity. On the other hand, runs using tour-

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

GECCO'16, July 20-24, 2016, Denver, Colorado, USA.

© 2016 ACM. ISBN TBA.

DOI: 10.1145/1235

nament selection show slow but steady increases in diversity, though not approaching that of lexicase selection runs [1].
GP parameters, etc.

3. RESULTS

SOOOO MANY GRAPHS!

3.1 Starting with high diversity

3.2 Starting with low diversity

3.3 Starting after a diversity crash

4. CONCLUSIONS

I'm hoping we have conclusions.

Acknowledgments

Lots of cool people helped us.

5. REFERENCES

- [1] T. Helmuth, N. F. McPhee, and L. Spector. Lexicase selection for program synthesis: a diversity analysis. In *Genetic Programming Theory and Practice XIII*, Genetic and Evolutionary Computation. Springer.
- [2] T. Helmuth and L. Spector. General program synthesis benchmark suite. In *GECCO '15: Proceedings of the 2015 Conference on Genetic and Evolutionary Computation*, July 2015.
- [3] T. Helmuth, L. Spector, and J. Matheson. Solving uncompromising problems with lexicase selection. *IEEE Transactions on Evolutionary Computation*, 19(5):630–643, Oct. 2015.
- [4] L. Spector. Assessment of problem modality by differential performance of lexicase selection in genetic programming: A preliminary report. In *1st workshop on Understanding Problems (GECCO-UP)*, pages 401–408, Philadelphia, Pennsylvania, USA, 7-11 July 2012. ACM.

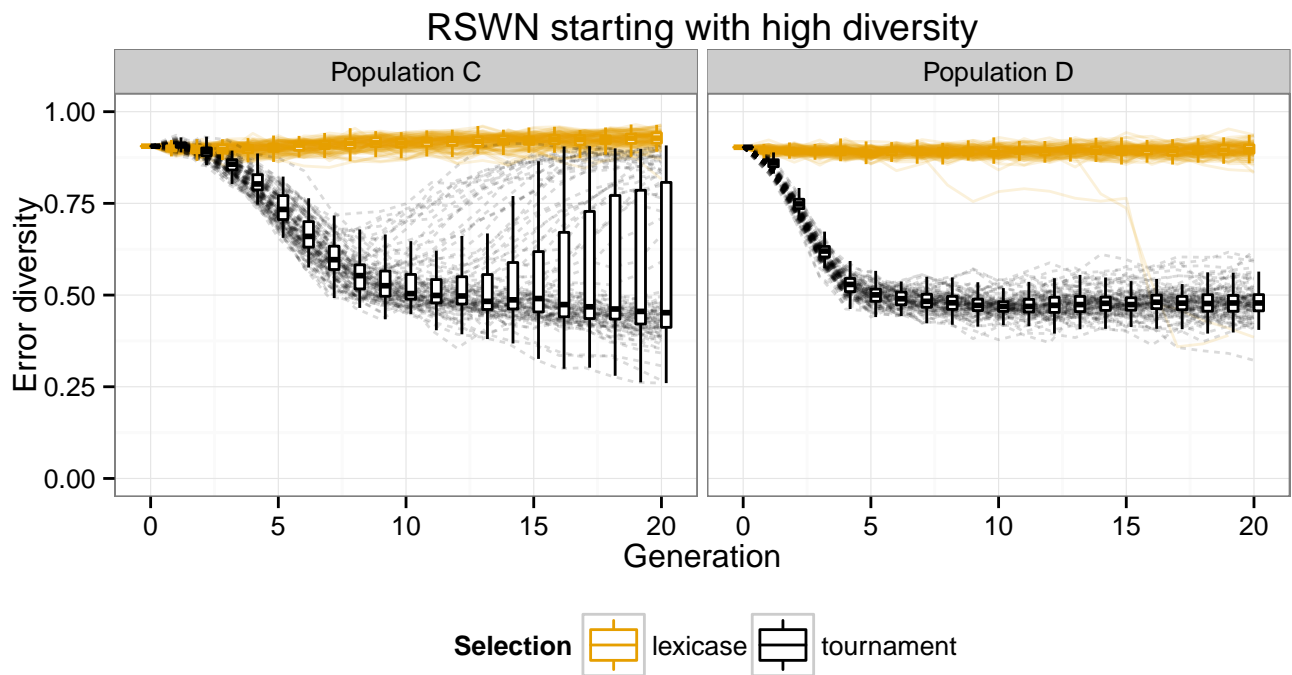


Figure 1: Changes in diversity over 100 “re-runs” of the replace-space-with-newline problem with both lexicase and tournament selections, starting from a population with high diversity coming from a lexicase selection run.

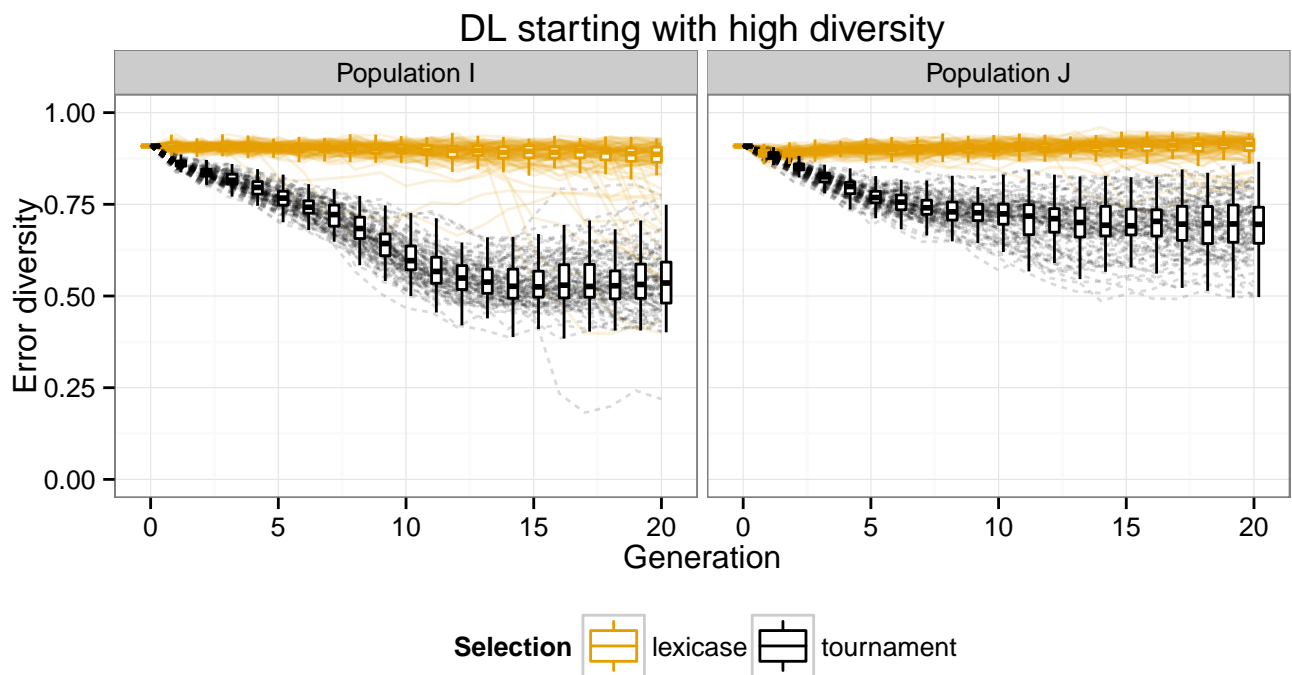


Figure 2: Changes in diversity over 100 “re-runs” of the double-letters problem with both lexicase and tournament selections, starting from a population with high diversity coming from a lexicase selection run.




Figure 3: Changes in diversity over 100 “re-runs” of the replace-space-with-newline problem with both lexicase and tournament selections, starting from a population with low diversity coming from a tournament selection run.

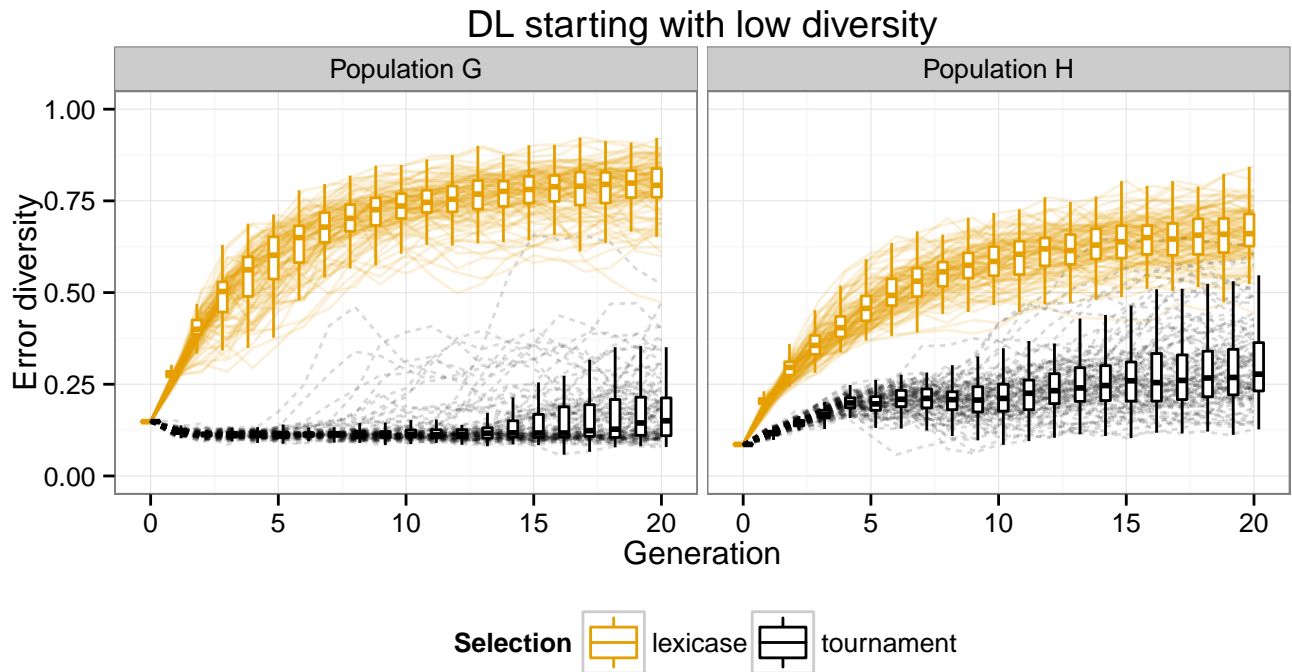


Figure 4: Changes in diversity over 100 “re-runs” of the double-letters problem with both lexicase and tournament selections, starting from a population with low diversity coming from a tournament selection run.

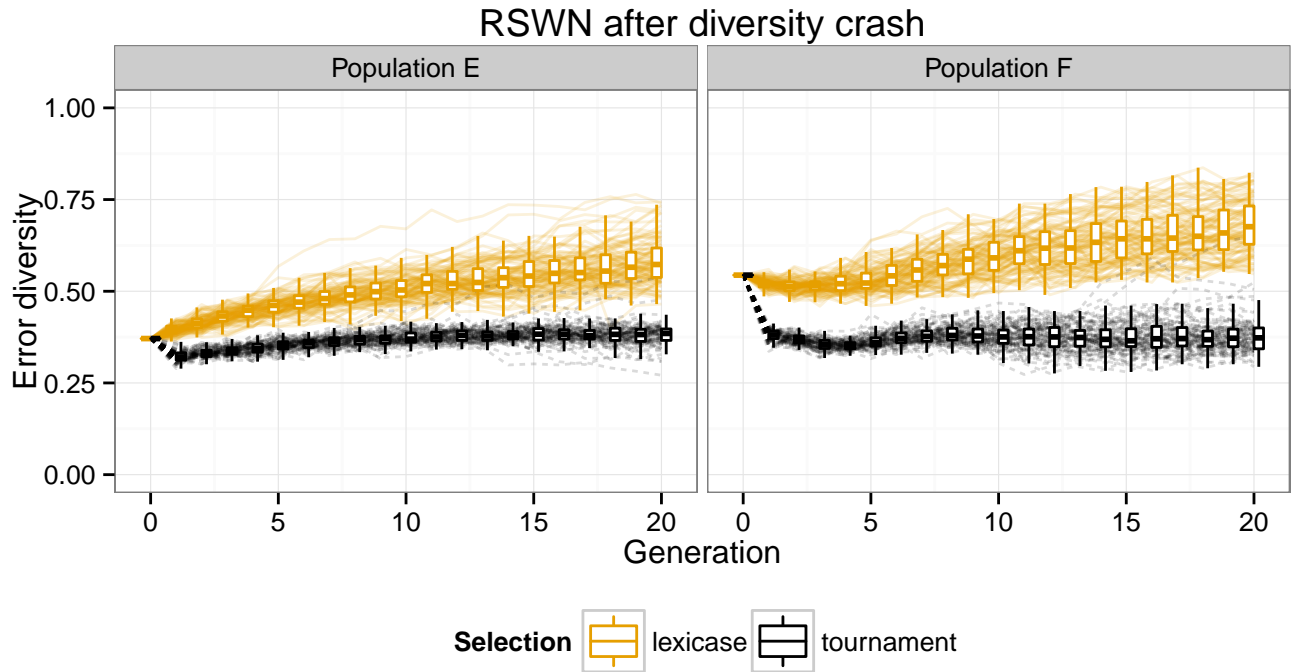


Figure 5: Changes in diversity over 100 “re-runs” of the replace-space-with-newline problem with both lexicase and tournament selections, starting from a population with that had lost diversity in a lexicase selection run.

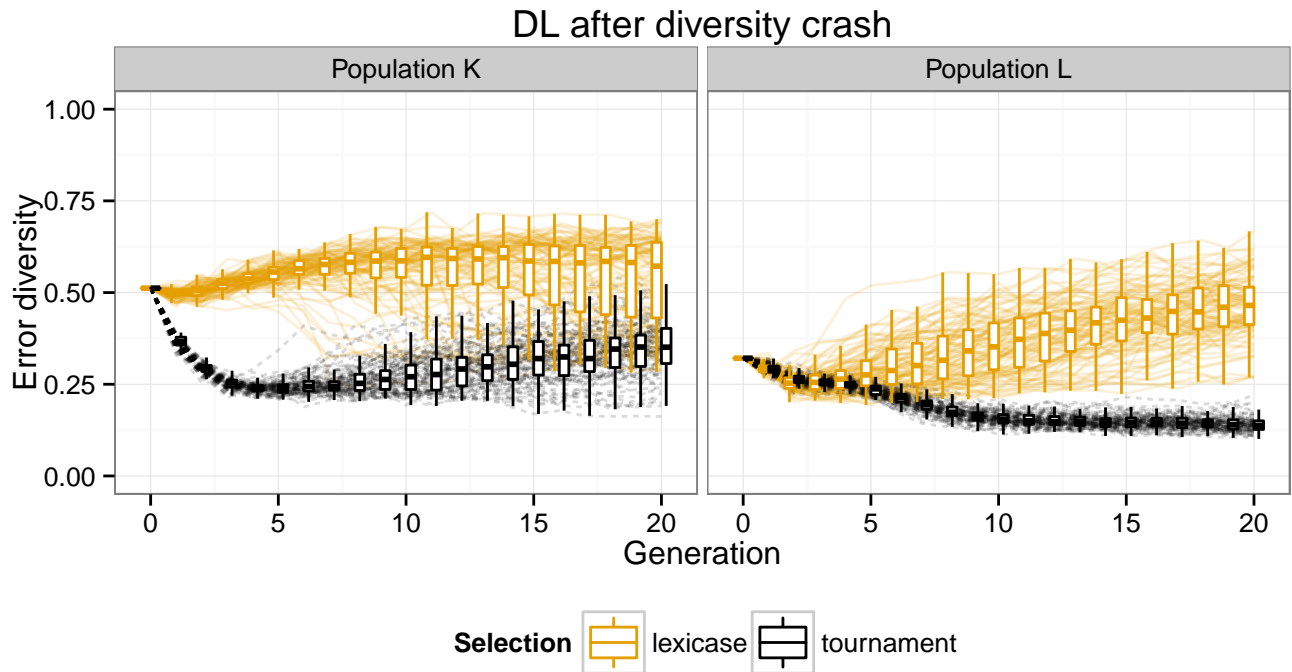


Figure 6: Changes in diversity over 100 “re-runs” of the double-letters problem with both lexicase and tournament selections, starting from a population with that had lost diversity in a lexicase selection run.