

Parent Selection and Diversification in Genetic Programming

Thomas Helmuth
Computer Science
Washington and Lee Univ
Lexington, Virginia
helmuth@wlu.edu

Nicholas Freitag McPhee
Div of Sci and Math
Univ of MN, Morris
Morris, MN 56267
mcphee@morris.umn.edu

Lee Spector
Cognitive Science
Hampshire College
Amherst, MA
lspector@hampshire.edu

ABSTRACT

More things!

Keywords

lexicase selection, hyperselection, PushGP, other stuff

1. INTRODUCTION

I bet we start here!

Lexicase selection [8] is nifty, eh?

Hyperselection provided initial motivation [3], but later we became interested more generally in what lexicase and tournament do differently starting with the same population.

[TMH: I'm not sure if we want to talk about error diversity in the Intro or somewhere else. But, the following paragraph could be moved to be the first paragraph of Experimental Design if we don't need to talk about diversity earlier]

In this paper we concentrate on diversity measures related to the outputs of the programs. One such diversity measure, *behavioral diversity*, has been shown to have correlation with problem-solving performance [6]. In behavioral diversity, the output of each program is recorded on each training case input and stored as a *behavior vector*. Behavioral diversity is then the percentage of distinct behavior vectors in the population. *Error diversity*, a slight variation of behavioral diversity, considers the percentage of distinct error vectors in the population where each error vector is computed by applying the error function to each output in the behavior vector. We believe error diversity does a good job of measuring how well evolution is exploring meaningful differences between programs that might be lost with a diversity measure that only takes into account syntactic (genotypic) diversity of the population, where two wildly different programs may actually compute the same function.

2. EXPERIMENTAL DESIGN

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

GECCO'16, July 20-24, 2016, Denver, Colorado, USA.

© 2016 ACM. ISBN TBA.

DOI: 10.1145/1235

Previous work has shown that using lexicase selection results in higher population error diversity than tournament selection across a variety of problems [2, 5]. These papers examined the diversity of entire GP runs, each starting with a different initial population and random number seed.

Here we examine the effects of these parent selection methods on population error diversity starting from specific population conditions besides a random initial population. In particular, we want to see how each method changes diversity in populations that occur naturally during an evolutionary run.

In order to produce the populations on which to experiment, we started GP runs and let them continue until they met certain stopping conditions; we then stored those populations and later conducted multiple trials with different random number seeds starting with those stored populations. We used three different stopping conditions in order to generate naturally occurring populations with interesting properties:

1. In a run using lexicase selection, we stopped if the population error diversity was greater than 0.9. This results in very diverse populations, allowing us to observe whether evolution is able to maintain such high diversity in the following generations.
2. In a run using tournament selection, we stopped if the population error diversity was less than 0.15. These populations allow us to see if methods promote diversification starting from such undiverse populations. They also allow us to see if methods perform differently on a population produced by tournament selection versus one produced by lexicase selection.
3. As described above, we were initially motivated here by observations of runs using lexicase selection that underwent major drops in diversity following hyperselection events, where one or a few individuals in the population received the majority of the parent selections in a generation. We had anecdotally noticed rapid diversity recovery following these events, but not examined them systematically [3].

In this condition, we stopped a run using lexicase selection when the error diversity reached a level at least 0.25 less than it had been at some point in the previous 10 generations. This allowed us to detect populations that had recently undergone large drops in diversity. We do not definitively know whether those drops are related to hyperselection events, but we expect that they are.

Table 1: PushGP parameters

Parameter	Value
runs per problem/parent selection combination	100
population size	1000
maximum genome size	1600
maximum initial genome size	400
Genetic Operator	Prob
alternation	0.2
uniform mutation	0.2
uniform close mutation	0.1
alternation followed by uniform mutation	0.5

In all three conditions, we only considered populations occurring after the first 10 generations in order to give evolution a chance to settle down after the extreme shifts that can happen at the beginning of a run.

In each trial, we continued running GP on a stored population for 20 generations and recorded the population error diversity. For each parent selection setting (lexicase and tournament selections), we conducted 100 trials with different random number seeds from each stored population.

We conducted these tests on two problems taken from a recent program synthesis benchmark suite [4]. The first problem, Replace Space With Newline (RSWN), searches for a program that takes as input a string and both prints the string after replacing all of the spaces in the input with newline characters and functionally returns the number of non-whitespace characters in the string. Previous examinations of error vector diversity on the RSWN problem indicate that lexicase selection maintains significantly higher diversity than tournament selection, which across 100 runs never achieved a median diversity higher than 0.25 [2].

The second problem, Double Letters, asks for a program that takes a string as input and prints the string after doubling every alphabetic character and tripling every exclamation point. All other characters should be printed once. As with the RSWN problem, lexicase selection consistently achieves high diversity on this problem. Differently than RSWN, runs using tournament selection show slow but steady increases in diversity, though not approaching that of lexicase selection runs [2].

For our experiments we used PushGP [11, 10], a stack-based genetic programming system.¹ PushGP supports a variety of control structures and multiple data types, making it a good choice for program synthesis tasks such as the problems we explore here. Except for parent selection, we used the exact same PushGP parameters in both the initial runs used to store interesting populations as well as the continuations of the stored populations. We give the most relevant parameters in Table 1. The parameters not listed here exactly follow those used in the experiments in [1].

These runs use the most recent version of PushGP, in which individuals are stored as linear genomes that we translate into hierarchical Push programs prior to execution [1]. These linear genomes admit a range of uniform genetic operators; we use four, listed in Table 1 with their probabilities. Alternation is a linear crossover operator modeled after the

sexual portion of ULTRA [9]. Uniform mutation may replace each instruction with 1% probability. Uniform close mutation may add or remove parentheses from the program. Finally, the last operator runs alternation on two parents and then uniform mutation on that child to produce a new child.

3. RESULTS

Using the techniques presented in the previous section we obtained populations on which to perform continuation experiments. For each combination of the 2 problems and 3 stopping conditions we stored populations from 2 runs, for a total of 12 populations. In the following subsections we group the results based on the stopping conditions, since they produce the most similar populations.

Starting with each stored population we conducted 100 “continuation” GP runs with lexicase selection and 100 with tournament selection. We let each continuation go for 20 generations, and plot the population error diversity across the runs. In particular, each figure has a standard box-and-whisker plot for each generation, with the box showing the median and quartiles. The whiskers stretch to the maximum and minimum values, ignoring outliers. In each figure we also plot the error diversity of each individual run at each generation, giving another way of visualizing the spread of diversities across runs and making it easier to trace runs with outliers.

Note that in a few settings, one or two runs out of 100 found solutions to the problem before the end of 20 generations. In these cases, we terminate the runs, and they do not contribute data past their termination generation. We do not believe these solutions have a large effect on the plots since no plot had more than two of these early-terminating runs.

3.1 Starting with high diversity

In this subsection we continue runs using stored populations evolved using lexicase selection that achieved error diversity greater than 0.9. As such, the initial populations of the runs have very high diversity, with most individuals producing distinct error vectors.

Figure 1 plots the continued runs started from two populations (C and D) stored from GP evolving on the RSWN problem. Lexicase selection consistently maintains high levels of diversity starting from both populations, with little variance. On the other hand, both plots show runs using tournament selection quickly losing significant diversity within the first 5 to 10 generations of the continuation, dropping from over 90% distinct error vectors down to around 50% distinct error vectors. Interestingly, the tournament selection runs on Population C show large differences in diversity the last 10 generations, with some becoming even less diverse while others recovering much of the lost diversity. On the other hand, the tournament selection runs on Population D maintain much more consistent diversity, with most runs having between 40% and 60% diversity in the remaining generations. It is unclear to us why these two populations result in such different diversity plots for tournament selection, but we assume it has to do with the composition of the stored population.

Figure 2, which plots the diversities of continuations of two populations (I and J) on the Double Letters problem, shows similar trends in both lexicase selection and tourna-

¹Lexicase selection has also been shown to be effective in tree-based genetic programming [5, 7].

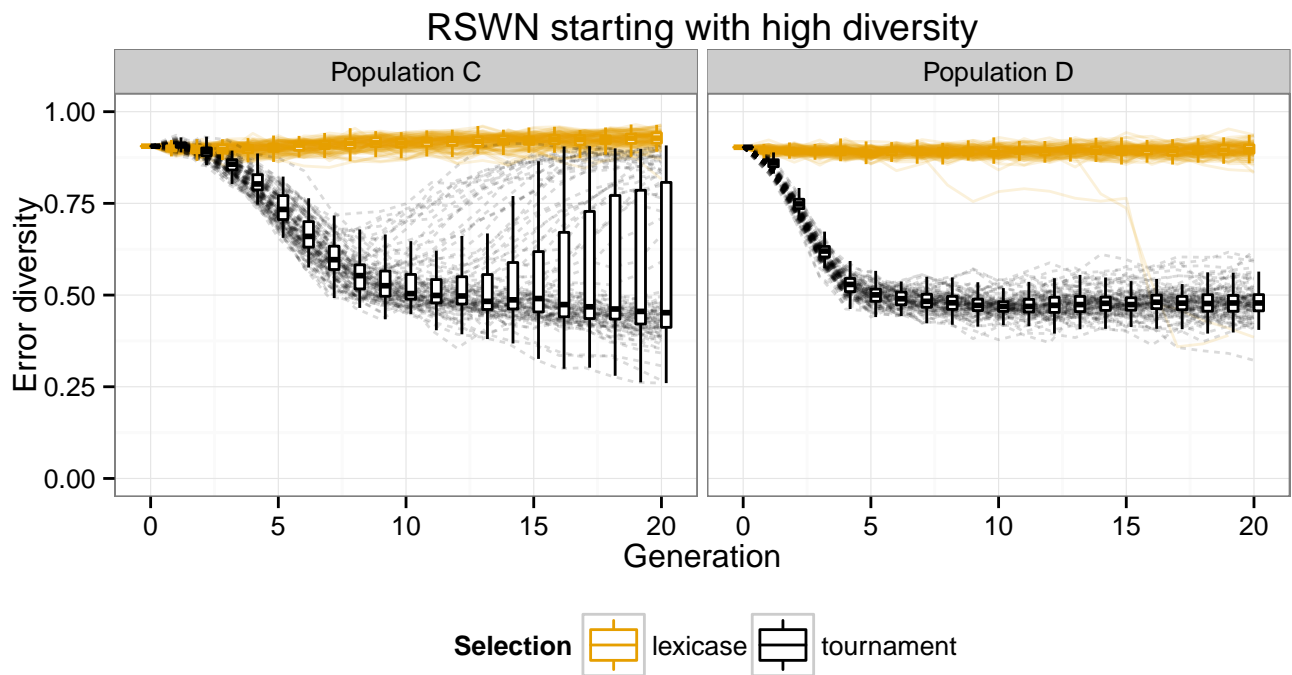


Figure 1: Error diversity over 100 continuations of the RSWN problem with both lexicase and tournament selections, starting from a population with high diversity naturally occurring in a run using lexicase selection.




Figure 2: Error diversity over 100 continuations of the double-letters problem with both lexicase and tournament selections, starting from a population with high diversity naturally occurring in a run using lexicase selection.

ment selection. Note that tournament selection maintains higher diversity on this problem than on the RSWN problem, though not as high as lexicase selection. This trend mirrors what has been previously seen on full GP runs [2].

[*should this paragraph be in Discussion?*] The continuations starting from high-diversity populations clearly show that lexicase selection can maintain a high population diversity while tournament selection cannot reliably do so.

One interesting observation in these plots is the occasional steep drop in diversity in a small number of runs using lexicase selection, which can be seen especially clearly in populations D and I. Based on similar runs we have encountered previously, we would guess that these runs underwent hyperselection events in which one or a small number of individuals were selected to make most of the children in a single generation. Hyperselection events can, understandably, lead to diversity crashes since most of the individuals in the population are closely related. Interestingly, previous work has shown that such events are neither a driving force or a hinderance in runs using lexicase selection [3].

3.2 Starting with low diversity

3.3 Starting after a diversity crash

4. DISCUSSION

Why does lexicase do so much better at diversification??? Does it have to do with specialists? Probably! Does it have to do with concentrating on individual test cases or combinations thereof? Probably!

Nic:

Some of the individual drops in lexicase diversity are likely due to hyperselection events. Some of those then recovered, but it's not clear that all did; not sure what we say about that part.

The details clearly depend on the problem and the run (i.e., the starting point). Not surprising, but definitely visible in the data.

5. CONCLUSIONS

I'm hoping we have conclusions.

Acknowledgments

Lots of cool people helped us.

6. REFERENCES

- [1] T. Helmuth. *General Program Synthesis from Examples Using Genetic Programming with Parent Selection Based on Random Lexicographic Orderings of Test Cases*. Ph.D. dissertation, 2015.
- [2] T. Helmuth, N. F. McPhee, and L. Spector. Lexicase selection for program synthesis: a diversity analysis. In *Genetic Programming Theory and Practice XIII*, Genetic and Evolutionary Computation. Springer.
- [3] T. Helmuth, N. F. McPhee, and L. Spector. The impact of hyperselection on lexicase selection. In *GECCO '16: Proceedings of the 2016 Conference on Genetic and Evolutionary Computation*, July 2016.
- [4] T. Helmuth and L. Spector. General program synthesis benchmark suite. In *GECCO '15: Proceedings of the 2015 Conference on Genetic and Evolutionary Computation*, July 2015.
- [5] T. Helmuth, L. Spector, and J. Matheson. Solving uncompromising problems with lexicase selection. *IEEE Transactions on Evolutionary Computation*, 19(5):630–643, Oct. 2015.
- [6] D. Jackson. Promoting phenotypic diversity in genetic programming. In *PPSN 2010 11th International Conference on Parallel Problem Solving From Nature*, volume 6239 of *Lecture Notes in Computer Science*, pages 472–481, Krakow, Poland, 11–15 Sept. 2010. Springer.
- [7] P. Liskowski, K. Krawiec, T. Helmuth, and L. Spector. Comparison of semantic-aware selection methods in genetic programming. In *GECCO 2015 workshop on Semantic Methods in Genetic Programming*. ACM, 2015.
- [8] L. Spector. Assessment of problem modality by differential performance of lexicase selection in genetic programming: A preliminary report. In *1st workshop on Understanding Problems (GECCO-UP)*, pages 401–408, Philadelphia, Pennsylvania, USA, 7–11 July 2012. ACM.
- [9] L. Spector and T. Helmuth. Uniform linear transformation with repair and alternation in genetic programming. In R. Riolo, J. H. Moore, and M. Kotanchek, editors, *Genetic Programming Theory and Practice XI*, Genetic and Evolutionary Computation, chapter 8, pages 137–153. Springer, Ann Arbor, USA, 9–11 May 2013.
- [10] L. Spector, J. Klein, and M. Keijzer. The Push3 execution stack and the evolution of control. In *GECCO 2005: Proceedings of the 2005 conference on Genetic and evolutionary computation*, pages 1689–1696, Washington DC, USA, 2005. ACM Press.
- [11] L. Spector and A. Robinson. Genetic programming and autoconstructive evolution with the push programming language. *Genetic Programming and Evolvable Machines*, 3(1):7–40, Mar. 2002.

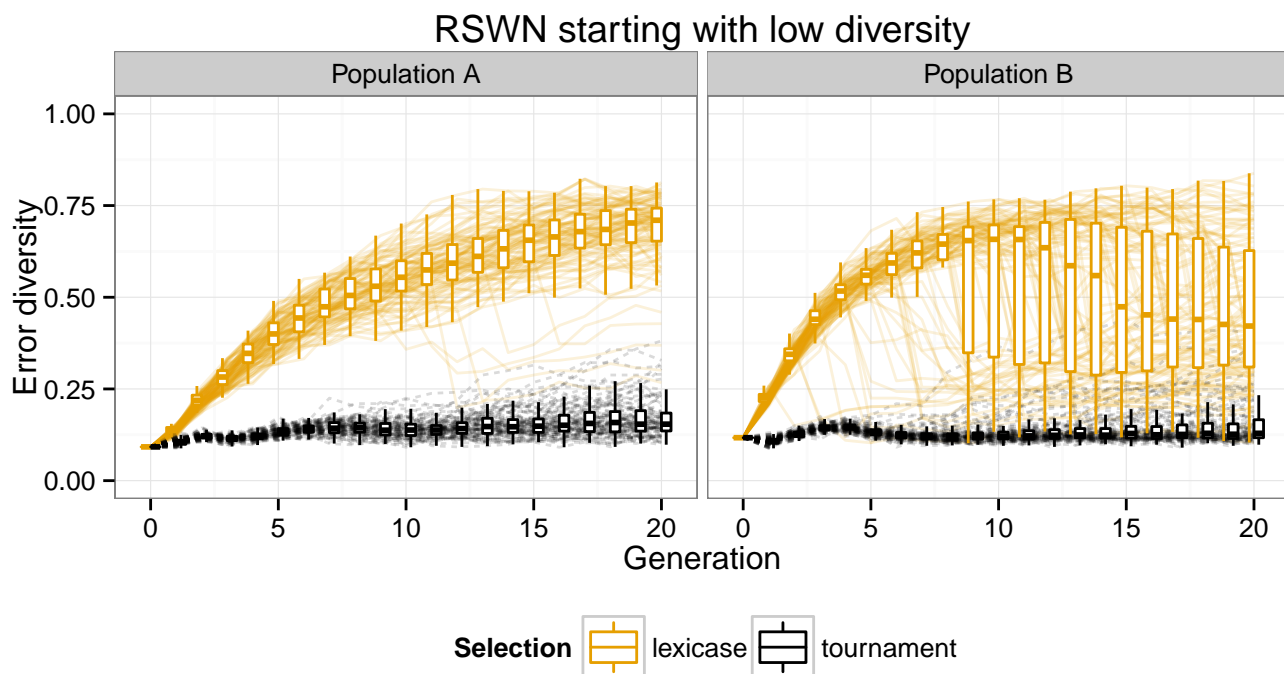


Figure 3: Error diversity over 100 continuations of the RSWN problem with both lexicase and tournament selections, starting from a population with low diversity naturally occurring in a run using tournament selection.




Figure 4: Error diversity over 100 continuations of the double-letters problem with both lexicase and tournament selections, starting from a population with low diversity naturally occurring in a run using tournament selection.




Figure 5: Error diversity over 100 continuations of the RSWN problem with both lexicase and tournament selections, starting from a population that had lost diversity in a diversity crash in a lexicase selection run.




Figure 6: Error diversity over 100 continuations of the double-letters problem with both lexicase and tournament selections, starting from a population that had lost diversity in a diversity crash in a lexicase selection run.